

# **Data 201**

## **Intro to Data Analysis in Python**

### **Data Wrangling with pandas for R Users**

#### **Week 2**

Dr. Rebin Muhammad

# Welcome – Week 2

Last week: Python fundamentals, NumPy, and the R→Python mindset.  
This week we focus on **tabular data**: the pandas DataFrame and how it maps from dplyr.

**Today** : DataFrame basics, dplyr→pandas verbs, method chaining, grouped summaries, missing data, and hands-on translation.

Have the **Week 2 Notebook** and your **R→Python cheat sheet** open.

# What we'll do today (Week 2)

1. **pandas in the ecosystem** — Where it fits (NumPy → pandas → seaborn/statsmodels)
2. **DataFrame basics** — Create, inspect, columns and rows
3. **dplyr → pandas** — select, filter, mutate, arrange (with correct boolean logic)
4. **Pipes → method chaining** — Build pipelines step by step
5. **group\_by and summarize** — Grouped summaries with `.groupby()` and `.agg()`
6. **Missing data** — Detect, drop, fill
7. **Practice** — Translate a dplyr pipeline into pandas

# Learning goals

By the end of today you will be able to:

1. **Explain** how pandas DataFrames relate to R data frames
2. **Perform** filtering, selecting, mutating, and sorting in pandas
3. **Use method chaining** instead of R pipes (`%>%`)
4. **Create grouped summaries** with `.groupby()` and `.agg()`

# pandas in the Python ecosystem

**Where pandas fits:**

- **NumPy** — numerical foundation (arrays, vectorization)
- **pandas** — tabular data: columns, rows, and dplyr-like operations
- **seaborn / statsmodels** — build on pandas DataFrames

**Think of pandas as:** R data frames + dplyr verbs. Same mental model; different syntax.

# DataFrame basics: R vs pandas

| Task              | R                            | pandas                         |
|-------------------|------------------------------|--------------------------------|
| Create            | <code>data.frame(...)</code> | <code>pd.DataFrame(...)</code> |
| Structure / types | <code>str(df)</code>         | <code>df.info()</code>         |
| First rows        | <code>head(df)</code>        | <code>df.head()</code>         |
| Dimensions        | <code>dim(df)</code>         | <code>df.shape</code>          |
| Column names      | <code>names(df)</code>       | <code>df.columns</code>        |

**Same idea:** a rectangular table with named columns and (usually) a row index.

→ Notebook Section 1

# Columns and rows

**Accessing columns:** one column → Series; multiple → DataFrame.

```
● ● ●  
1 df[ "price" ]           # one column → Series  
2 df[ [ "price" , "size" ] ] # multiple columns → DataFrame
```

**Filtering rows:** pass a boolean Series (same length as the DataFrame).

```
● ● ●  
1 df[df[ "price" ] > 300]
```

→ Notebook Section 2

# Boolean logic (critical for filters)

In R you use & and | ; pandas is the same, but **you must use parentheses** so each condition is a complete expression.

| R                          | pandas                         |
|----------------------------|--------------------------------|
| & and                      | & and                          |
| Vector recycling           | No recycling (aligns by index) |
| Parentheses often optional | <b>Parentheses required</b>    |

**Wrong:** `df[df["price"] > 300 & df["size"] > 1000]` → operator precedence error.

**Right:** `df[(df["price"] > 300) & (df["size"] > 1000)]`

# Boolean logic – example

```
1 # One condition
2 df[df["price"] > 300]
3
4 # Two conditions (AND): use parentheses around each comparison
5 df[(df["price"] > 300) & (df["size"] > 1000)]
6
7 # OR
8 df[(df["price"] > 500) | (df["size"] < 500)]
```

**Rule of thumb:** wrap each comparison in parentheses when combining with `&` or `|`.

# `select()` → column subsetting

| dplyr                     | pandas                      |
|---------------------------|-----------------------------|
| <code>select(a, b)</code> | <code>df[["a", "b"]]</code> |

- **No tidyselect helpers by default** — use explicit column names (or a list of names).
- Single column: `df["a"]` (Series). Multiple: `df[["a", "b"]]` (DataFrame).



```
1 df[["price", "size", "neighborhood"]]
```

→ Cheat Sheet Section 6

# filter() → row filtering

**Option 1:** Boolean indexing (like R's logical subsetting).



```
1 df[df["price"] > 200]
```

**Option 2:** .query( ) — often feels closest to dplyr.



```
1 df.query("price > 200")
2 df.query("price > 200 and size > 1000")
```

**Note:** .query( ) uses string expressions; column names must be valid Python identifiers (or use backticks in the string for special names).

→ Notebook Section 2

# **mutate() → .assign()**

**dplyr**

```
mutate(ppsqft = price / size)
```

**pandas**

```
.assign(ppsqft=lambda d:  
       d.price / d.size)
```

**Why lambda?** The argument is the (current) DataFrame. Using `d` avoids confusion with partially-created columns and makes each new column depend only on existing columns.



```
1 df.assign(ppsqft=lambda d: d.price / d.size)
```

**Important:** `.assign( )` returns a *new* DataFrame; it does not modify the original (immutable style).

→ Notebook Section 3

# arrange() → .sort\_values()

| dplyr                | pandas                                    |
|----------------------|---|
| arrange(price)       | .sort_values("price")                     |
| arrange(desc(price)) | .sort_values("price",<br>ascending=False) |
| Multiple columns     | .sort_values(["col1", "col2"])            |



```
1 df.sort_values("price")
2 df.sort_values("price", ascending=False)
3 df.sort_values(["neighborhood", "price"])
```

# Pipes → method chaining

**R:** `%>%` passes the result of the left side as the first argument to the right.

**Python:** Each method returns a DataFrame, so you chain with dots. Use parentheses so you can break the chain across lines.

**R:**

```
● ● ●  
1 df %>%  
2   filter(price > 200) %>%  
3   mutate(ppsqft = price / size)
```

**Python:**

```
● ● ●  
1 (df  
2   .query("price > 200")  
3   .assign(ppsqft=lambda d: d.price / d.size))
```

# Why method chaining helps

- **Readable:** Top to bottom = order of operations (filter → mutate → ...).
- **No temporary variables:** Each step is the input to the next.
- **Same as dplyr:** If you're used to pipes, chaining will feel familiar.

**Style:** Put the opening ( on the same line as the object, then each method on its own line with a leading dot.

# group\_by() → .groupby()

| dplyr          | pandas        |
|----------------|---------------|
| group_by(x)    | .groupby("x") |
| summarize(...) | .agg(...)     |

**Typical pattern:** .groupby("col").agg(new\_name= ("existing\_col", "function"))



```
1 df.groupby("neighborhood").agg(mean_price=( "price", "mean"))
2 df.groupby("type").agg(
3     mean_price=( "price", "mean"),
4     count=( "price", "count")
5 )
```

→ Notebook Section 4

# Grouped summary – full example

R:

```
● ● ●  
1 df %>%  
2   group_by(neighborhood) %>%  
3   summarize(mean_price = mean(price), n = n())
```

Python:

```
● ● ●  
1 df.groupby("neighborhood").agg(  
2     mean_price=("price", "mean") ,  
3     n= ("price", "count")  
4 )
```

**Note:** pandas uses "count" or "size" for row counts; here ("price", "count") counts non-null price values.

# Missing data in pandas

| Task   | R                                | pandas                              |
|--------|----------------------------------|-------------------------------------|
| Detect | <code>is.na(x)</code>            | <code>df.isna() or pd.isna()</code> |
| Drop   | <code>na.omit()</code>           | <code>.dropna()</code>              |
| Fill   | <code>replace(), ifelse()</code> | <code>.fillna()</code>              |



```
1 df.isna()           # DataFrame of True/False
2 df.dropna()         # drop rows with any NA
3 df.dropna(subset=["price"]) # drop only where price is NA
4 df.fillna(0)        # fill NAs with 0
```

→ Cheat Sheet Section 7

# Series vs DataFrame (a common pitfall)

- **One column** `df[ "price" ]` → **Series** (1D). Result of `.agg( "mean" )` on a single column can be a Series.
- **Two or more columns** `df[ [ "price", "size" ] ]` → **DataFrame** (2D).

**Why it matters:** Some operations expect a DataFrame (e.g. chaining `.assign( )`). If you have a Series, you may need to wrap or use `.to_frame( )`.

**Rule of thumb:** Subset columns with `df[ [ "a", "b" ] ]` when you want to keep a DataFrame.

# Common pandas pitfalls

- **Forgetting parentheses in filters** — `(df["price"] > 300) & (df["size"] > 1000)` needs parentheses around each comparison.
- **Confusing Series vs DataFrame** — Single column is a Series; use `df[ [ "col" ] ]` for a one-column DataFrame.
- **Modifying views vs copies** — Prefer `.assign()` and new objects; avoid chained indexing (e.g. `df[df.x > 0]["y"] = 1`).
- **Overusing loops** — Use vectorized operations and `.apply()` when appropriate; avoid row-by-row loops on large data.

→ Notebook Section 5

# Active learning (Part 1) – 20-25 min

Translate this dplyr pipeline to pandas.

R code:

```
● ● ●  
1 df %>%  
2   filter(price > 300, size > 1000) %>%  
3   mutate(ppsqft = price / size) %>%  
4   group_by(type) %>%  
5   summarize(mean_ppsqft = mean(ppsqft))
```

Tasks:

1. **Predict** — What is the structure of the result? (One row per what? Which columns?)
2. **Translate** — Write the pandas version (.query( ) or boolean filter, .assign( ), .groupby( ), .agg( )).
3. **Run** — Execute in the Week 2 Notebook and interpret.

# Active learning – solution sketch

```
● ● ●  
1 (df  
2 .query("price > 300 and size > 1000")  
3 .assign(ppsqft=lambda d: d.price / d.size)  
4 .groupby("type")  
5 .agg(mean_ppsqft=("ppsqft", "mean"))  
6 )
```

**Result:** One row per type; column mean\_ppsqft = mean of price/size in each group.

Compare with your solution. Questions on .query( ), .assign( ), or .agg( )?

# Active learning (Part 2) – if time, 10–15 min

**Extend the pipeline:** From the same `df`, add a step that keeps only neighborhoods with at least 5 rows, then compute mean price by neighborhood.

**Hints:** `.groupby().filter()` or `.groupby().agg()` then filter the result; or use `.groupby().agg()` with a count column and then subset rows.

- Use the Week 2 Notebook; compare with a neighbor or the solution.

# Active learning (Part 2) –



```
1 # Solution A (recommended): keep groups with at least 5 rows, then mean price
2 result = (
3     df
4     .groupby("neighborhood")
5     .filter(lambda g: len(g) >= 5)
6     .groupby("neighborhood")["price"]
7     .mean()
8 )
9 result
```



```
1 # Solution B: summarize count + mean, then keep count >= 5
2 summary = (
3     df
4     .groupby("neighborhood")
5     .agg(n=("price", "count"), mean_price=("price", "mean"))
6 )
7
8 summary[summary["n"] >= 5][["mean_price"]]
```

# Key takeaways

1. **pandas mirrors dplyr conceptually** — select, filter, mutate, arrange, group\_by, summarize all have direct equivalents.
2. **Syntax is more explicit** — Column names in quotes; boolean filters need parentheses; new columns via `.assign(..., lambda d: ...)`.
3. **Method chaining replaces pipes** — `(df .query(...).assign(...).groupby(...).agg(...))` reads like a dplyr pipeline.
4. **Series vs DataFrame** — Know when you have one column (Series) vs a table (DataFrame); use `df[["col"]]` when you need a DataFrame.

# Looking ahead

## **Next week (Week 2-3): Visualization in Python**

- ggplot2 → seaborn
- Grammar-of-graphics thinking (data, aesthetics, geoms) mostly survives
- Bring your R→Python cheat sheet

Week 2's wrangling skills are what you'll feed into plots and models in the rest of the course.

# Thank you

**Data 201 · Intro to Data Analysis in Python**

**Week 2-Part 1 — Data Wrangling with pandas for R Users**

